CSC411H1S Project 4

Hao Hui Tan(999741711, tanstev1)
Kyle Zhou (1000959732, zhoukyle)

March 29, 2018

1. The board is represented by a flat 9-element NumPy tuple. Turn denotes whose turn it is (1 for X, 2 for O). Done denotes whether the game is done (True if game is over, False otherwise.)

   Below is an example of a sample game played against myself.

   ```
   Python 2.7.14 |Anaconda custom (64-bit)| (default, Oct  5 2017, 02:28:52)
   [GCC 4.2.1 Compatible Clang 4.0.1 (tags/RELEASE_401/final)] on darwin
   env.render()
   ====
   env.step(0)
   Out[3]: (array([1, 0, 0, 0, 0, 0, 0, 0, 0]), 'valid', False)
   env.render()
   x..
   ====
   env.step(4)
   Out[5]: (array([1, 0, 0, 0, 2, 0, 0, 0, 0]), 'valid', False)
   env.render()
   x..
   .o.
   ====
   env.step(8)
   Out[7]: (array([1, 0, 0, 0, 2, 0, 0, 0, 1]), 'valid', False)
   env.render()
   x..
   .o.
   ..x
   ====
   env.step(2)
   Out[9]: (array([1, 0, 2, 0, 2, 0, 0, 0, 1]), 'valid', False)
   env.render()
   x.o
   .o.
   ..x
   ====
   env.step(6)
   Out[11]: (array([1, 0, 2, 0, 2, 0, 1, 0, 1]), 'valid', False)
   env.render()
   x.o
   .o.
   x.x
   ====
   env.step(3)
   Out[13]: (array([1, 0, 2, 2, 2, 0, 1, 0, 1]), 'valid', False)
   env.render()
   x.o
   oo.
   x.x
   ====
   env.step(7)
   Out[15]: (array([1, 0, 2, 2, 2, 0, 1, 1, 1]), 'win', True)
   env.render()
   x.o
   oo.
   xxx
   ====
   ```

```
env.done
Out[17]: True
env.step(1)
Out[18]: (array([1, 0, 2, 2, 2, 0, 1, 1, 1]), 'done', True)
```

2. (a) The following is the new implemented policy

Listing 1:

```
1  class Policy(nn.Module):
2      """
3      The Tic-Tac-Toe Policy
4      """
5      def __init__(self, input_size=27, hidden_size=64, output_size=9):
6          super(Policy, self).__init__()
7
8          self.linear1 = nn.Linear(input_size, hidden_size)
9          self.linear2 = nn.Linear(hidden_size, output_size)
10
11      def forward(self, x):
12          x = F.relu(self.linear1(x))
13          x = self.linear2(x)
14          return F.log_softmax(x)
```

(b) The 27 dimensions are a flattened encoding of a one-hot encoding of the state of the board. If `.view(3,9)` is applied to the array, the columns would be the one-hot encoding of each cell in the board (starting from the top left, going across each row, and ending in the bottom right).

If a column contains "1 0 0," the cell is empty.
If a column contains "0 1 0," the cell is occupied by an X.
If a column contains "0 0 1," the cell is occupied by an O.

(c) The value in each dimension means the chance that making the move (e.g. adding an X into that cell) would result in winning the game. This policy is stochastic because it samples the next move from a distribution, rather than following a deterministic algorithm.

3. (a) The implementation of **compute_returns** is shown below.

Listing 2:

```
1  def compute_returns(rewards, gamma=1.0):
2      """
3      Compute returns for each time step, given the rewards
4        @param rewards: list of floats, where rewards[t] is the reward
5                          obtained at time step t
6        @param gamma: the discount factor
7        @returns list of floats representing the episode's returns
8            G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + ...
9
10     >>> compute_returns([0,0,0,1], 1.0)
11     [1.0, 1.0, 1.0, 1.0]
12     >>> compute_returns([0,0,0,1], 0.9)
13     [0.7290000000000001, 0.81, 0.9, 1.0]
14     >>> compute_returns([0,-0.5,5,0.5,-10], 0.9)
15     [-2.5965000000000003, -2.8850000000000002, -2.6500000000000004,
            -8.5, -10.0]
16     """
17     # TODO
```

```
18        result = []
19        for index in range(len(rewards)):
20            sum_returns = 0
21            power = 0
22            for i in range(index, len(rewards)):
23                sum_returns = sum_returns + ((gamma ** power) * rewards[i
                    ])
24                power = power + 1
25            result.append(sum_returns)
26        return result
```

(b) If we update the weights in the middle of an episode, it would lead to erroneous results since the algorithm does not know the terminal state of the episode when it's updating the weights. For instance, if the terminal state for an agent is "loss", then this will cancel out all the rewards that are associating with the agent's previous actions in the episode. Therefore, we should not update the weights in the middle of an episode.

Listing 3:

4. (a)
```
 1  def get_reward(status):
 2      """Returns a numeric given an environment status."""
 3      return {
 4          Environment.STATUS_VALID_MOVE   : 0,
 5          Environment.STATUS_INVALID_MOVE : -99,
 6          Environment.STATUS_WIN          : 1,
 7          Environment.STATUS_TIE          : 0,
 8          Environment.STATUS_LOSE         : -1
 9      }[status]
```

(b) The reward for a valid move is 0, since a single valid move doesn't directly contribute to a win or a loss. The reward for an invalid move is -99, since we really don't want our neural network to try invalid moves. Making invalid moves is worse than losing the game. The reward for a win is 2, and that's because we want to encourage our neural network to optimize for wins. The reward for a tie is 1, and that is because we want to encourage the neural network to finish the game, but we don't want the neural network to choose this option over a win. The reward for a loss is -1, since we want to discourage losses. The magnitude of these rewards are mostly close to zero, except for invalid moves. This is because all of the regular states are states that could occur in regular play, and since we really want to discourage invalid moves, having the rest of the rewards close together makes the -99 penalty more meaningful.

5. (a)

   (b)

   (c)

   (d)

6.

7.

8.